

Concordia University

Sepehr Seifpour 40197899

Safi Ullah 40222120

Bharat Saini 40202642

COMP 6721 Applied Artificial Intelligence

Professor: Dr. René Witte

Dataset

We have selected our dataset from the Kaggle, containing a diverse set of images of faces. The images are categorized into one of the seven emotions: Angry, Disgust, Fear, Happy, Sad, Surprise, and Neutral, based on the depicted facial expressions. The cleaning process of the dataset has been done to generate the data appropriate for our project. The selected dataset is divided into two sets for training and testing.

Dataset details on Kaggle:

- Dataset url : <https://www.kaggle.com/datasets/msambare/fer2013>
- Owner: Manas Sambare
- Authors: nil
- Coverage: nil
- DOI Citation: nil
- Provenance: nil
- License: nil

Train Dataset

Category	Number of Images	characteristics of the dataset
Bored	511	Frontal face shot
angry	552	Frontal face shot
neutral	500	Frontal face shot
focused	378	Frontal face shot

Number of Images for Training : 1911

Test Dataset

Category	Number of Images	characteristics of the dataset
Bored	138	Frontal face shot
angry	260	Frontal face shot

neutral	300	Frontal face shot
focused	381	Frontal face shot

Total Number of Images for testing: 1079

Total Number of Images: 2990

Justification for dataset choices:

The reason of choosing this dataset was basically because they were already labeled and it contains good range of categories to choose image from .for example during our searching for “bored” categories we face a really huge challenge to search and find a good dataset but we couldn't find any. Even though we found a huge dataset that was gathered from video students attending an online class and it was categorized with bored as they students were not engaged in the session but it was pretty useless as there were so many duplicate images and frames that could lead to confusion in the following steps. So we come up with solution of finding a really huge dataset that contains various and different category of face expression ,thus it makes it really easy for us to choose our images from it. As an example ,for “bored” category we realized that in our last huge dataset that we found there was missing match in “sad” folder as we could recognized that bored images were mixed with sad images so we easily cut those images and create out one “bored” category

Moreover it has angry prelabeled dataset ,although we need to go through the images and find ones with the same size and color in order to clean and categorized dataset

Also for focused category we decided to follow the same process by going through neutral category and find those images that have striking resemblance to focused category

One of the biggest advantage of our dataset is that it mostly contains more than 3000 in each category so it allowed to find and search the ones that we really need

Provenance Information:

Finding a good dataset was pretty time consuming for our team, as a matter of fact we reach to some resource in github that was private dataset and we need to email them to give us the access for using their dataset and still we finally realized that their dataset was useless because it mostly contains the different facial expression of few people that makes some boundaries for us in further step especially when we want to train our model with huge and different range of dataset

This is the different references the we tried to use their dataset but they were somehow useless for us:

https://github.com/engagenet/engagenet_baselines

And this is our main dataset that we are going to use in our project:

<https://www.kaggle.com/datasets/msambare/fer2013>

Data Cleaning

The dataset is processed and the cleaning process is implemented to ensure that the dataset satisfies our project requirements.

The data cleaning process involved the following key steps:

- **Image Quality Assessment:**
 - Images were inspected for issues such as low brightness, blurriness and low resolution. This ensured that the facial features were well-defined. Any images that did not meet the quality standards were excluded from the dataset.
 - Brightness: The images were examined to identify the images with low brightness. This ensured that the images in the dataset have sufficient brightness to recognize the facial expressions accurately.
 - Blurriness: The images are inspected to find any blurriness that could hinder the accuracy and performance of the model.
 - Low Resolution: Each image is examined to make sure that no image with low resolution is present in the dataset. This ensures that the model can read the facial features from the images clearly.
- **Noise reduction:**
 - Presence of noise in the images can interfere with the accuracy of the facial recognition. Techniques such as noise filtering were applied to identify and remove any unwanted data from the dataset.
- **Standardization of Image Format:**
 - All images were converted to consistent format (.jpg) and resolution. This helps us to maintain uniform data across the dataset. The dataset was thoroughly examined and consistency checks were run after performing standardization of the images to make sure that the dataset is consistent.
- **Data Augmentation:**
 - Techniques such as rotating, zooming, and flipping were used to enhance the dataset quality.
 - Rotating: Images were rotated to provide different angles to the faces. This ensures that the model is able to recognize emotions from different angles.

- Zooming: Images were zoomed to change the scale of facial features. This enables the model to adapt to varying differences between the subject and the camera. Hence, improving the performance of the model.
- Flipping: The images were flipped to generate mirror images. With the addition of these flipped images, the model will be able to recognize both sides of the faces, thereby increasing the accuracy of the model.
- **Annotation Validation:**
 - The existing dataset was verified to find and fix any discrepancies that might have arisen during the labelling of the original dataset.
 - Label consistency check: The images were examined to cross verify with the labels and the facial expression in the image. Any images where the labels did not match the facial expressions were corrected.
 - Ambiguous cases: The images that could satisfy multiple labels were examined carefully to find the appropriate label.

Labeling

To satisfy our requirements, we have performed a relabeling and filtering process on the original dataset. Emotions Categories:

- Angry
- Bored
- Focused
- Neutral

We partitioned and relabeled the data initially categorized as Sad to extract the Bored category, and similarly for Neutral to obtain the Focused category.

The final dataset details after relabelling are as follows:

- Total number of images: 3020
- Number of directories: 2 (Test and Train)

Number of images per class in test directory:

Angry	Bored	Focused	Neutral
260	138	381	300

Number of images per class in train directory:

Angry	Bored	Focused	Neutral
552	511	378	500

In order to have functional labeling we used a little but functional coding so it makes it really easy and fast for us to choose our dataset and then by going through them we would realize if it needs additional process or not.

```
In [11]: import os
import shutil
import random
# Define folder paths
neutral1_folder = "neutral_train/neutral"
neutral2_folder = "neutral_test/neutral"
neutral3_folder = "neutral_train_new"
neutral4_folder = "neutral_test_new"

focused_train_folder="Focused/Focused_Train"
focused_test_folder="Focused_Test"
bored_folder="Bored/Bored"
bored_train_new_folder="Bored_new"
sad_folder="sad"
# List files in neutral2 folder

myfunction(neutral2_folder,focused_test_folder,neutral4_folder,300)

In [3]: def myfunction(mainfilepath,comparefilepath,newfilepath,number):
allImages = os.listdir(mainfilepath)
selectedImages=os.listdir(comparefilepath)
non_selectedImages=[image for image in allImages if image not in selectedImages]
random_image=random.sample(non_selectedImages,number)
for image in random_image:
source_path = os.path.join(mainfilepath, image)
destination_path = os.path.join(newfilepath, image)
shutil.copy(source_path, destination_path)
```

This relabeled dataset now satisfies the requirements for our facial emotion recognition project.


Dataset Visualization

1-Class distribution

For plotting the train and test class distribution of their dataset , we actually create a function with folder name as a parameter , so at the end it plots the number of dataset for each categories

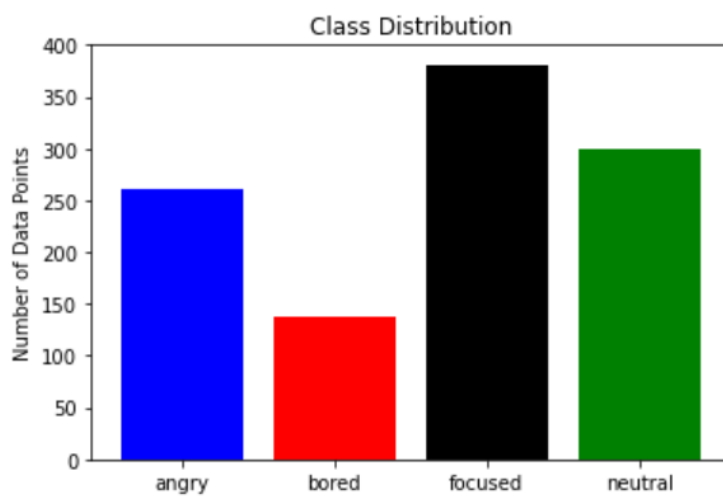
2-Train dataset:

```
visualize_class_distribution("train")
```



3-Test dataset :

```
visualize_class_distribution("test")
```



4-Sample Images

The given code selects a certain number of photographs at random and arranges them in a 5x5 grid. It does this by listing all image files inside a folder with popular image file extensions (jpg, jpeg, and png). Additionally, it resizes the photographs that are presented to the appropriate sizes, making it simple to view a portion of the folder's images in a grid that is organized.

```
def display_random_images_from_class_folder(class_folder, num_images=25, image_size=(128, 128)):

    # List all image files in the class folder
    image_files = [f for f in os.listdir(class_folder) if f.lower().endswith(('.jpg', '.jpeg', '.png'))]

    # Check if there are enough images to sample
    if num_images > len(image_files):
        num_images = len(image_files)

    # Randomly select the specified number of images
    selected_images = random.sample(image_files, num_images)


    num_rows, num_cols = 5, 5
    fig, axes = plt.subplots(num_rows, num_cols, figsize=(12, 12))

    for i, ax in enumerate(axes.ravel()):
        if i < num_images:
            image_file = selected_images[i]
            image_path = os.path.join(class_folder, image_file)
            img = Image.open(image_path)

            # Resize the image to the desired size
            img = img.resize(image_size)

            ax.imshow(img)
            ax.axis('off')

    plt.show()
```



5-Pixel Intensity distribution:

The given code first gathers all of the picture files inside of a folder that ends in .jpg. It then chooses a predetermined number of these files at random and shows their pixel intensity histograms. Through the organized display of a sample of photos and the related pixel intensity distributions, this visual representation facilitates comprehension of the image collection by offering insights into the content and pixel properties of the dataset.:

```
def visualize_sample_images_and_pixel_intensity(folder_path, num_samples=5):
    # Get a list of image file paths in the folder
    image_files = [os.path.join(folder_path, filename) for filename in os.listdir(folder_path) if filename.lower().endswith(('jpg', 'jpeg', 'png'))]

    # Randomly sample 'num_samples' images from the folder
    sample_images = random.sample(image_files, min(num_samples, len(image_files)))


    # Plot the images and their pixel intensity distributions
    fig, axes = plt.subplots(num_samples, 2, figsize=(12, 4 * num_samples))

    for i, image_path in enumerate(sample_images):
        # Display the image
        img = Image.open(image_path)
        axes[i, 0].imshow(img)
        axes[i, 0].set_title(f"Sample Image {i + 1}")
        axes[i, 0].axis('off')


        # Display the pixel intensity histogram
        pixel_values = np.array(img.convert('L')).ravel() # Convert to grayscale and flatten using NumPy
        axes[i, 1].hist(pixel_values, bins=256, range=(0, 256), density=True, color='blue', alpha=0.7)
        axes[i, 1].set_title("Pixel Intensity Histogram")

    plt.tight_layout()
    plt.show()
```

Sample Image 1



Pixel Intensity Histogram



CNN Architecture

1. Model Overview and Architecture Details:

The main model, SimpleCNN, is a convolutional neural network designed for analyzing facial images to detect emotions. It consists of three convolutional layers followed by max-pooling layers to capture hierarchical features.

The architecture details for the SimpleCNN model are as follows:

- Conv1: Convolutional layer with 1 input channel, 16 output channels, kernel size 5x5, and ReLU activation function.
- Conv2: Convolutional layer with 16 input channels, 32 output channels, kernel size 5x5, and ReLU activation function.
- Pooling: Max-pooling later with a kernel size of 2x2.
- Fully Connected Layers: After flattening the output, two fully connected layers follow with 32 * 6 * 6 and 128 neurons, respectively.
- Output Later: The final layer consists of four neurons corresponding to the four emotion classes.
- Dropout Layer: Additionally, the model includes a dropout later with a dropout rate of 0.5 for regularization.
- ReLU activation after each convolutional layer and the final layer.
- The SimpleCNN model has two variants, namely, SuperSimpleCNN and ImprovedCNN.

SuperSimpleCNN:

- This variant is a simplified version of the SimpleCNN model with changes in architecture.
- It includes two convolutional layers, conv1 and conv2 with LeakyReLU activation and batch normalization.
- LeakyReLU activation and batch normalization are applied before the fully connected layer.
- The model uses max-pooling after each convolutional layer.
- It has a single fully connected layer.

ImprovedCNN:

- It uses an improved convolutional neural network with three convolutional layers (conv1, conv2 and conv3).
- ReLU activation after each convolutional layer and the first fully connected layer.
- Batch normalization applied after the first two convolutional layers.
- Dropout introduced before the final fully connected layer for regularization.

2. Training Process:

- **Number of Epochs:**

SimpleCNN	SuperSimpleCNN	ImprovedCNN
50	45	58

- **Learning Rate:**

Learning rate is a crucial hyperparameter that influences the step size at which the optimizer updates the model parameters during training. Learning Rate is set to 0.0001 for all models.

```
optimizer = optim.Adam(model.parameters(), lr=0.0001)
```

- **Lost Function:**

The loss function measures the difference between the model's predictions and the actual labels. Cross-entropy loss function is used with class weights for imbalanced classes. It provides a mechanism to handle imbalanced class distributions during the training of a neural network.

```
criterion = nn.CrossEntropyLoss(weight=class_weights)
```

- **Optimization Algorithm:**

Adam optimizer is used for training. Adam is an optimization algorithm that is designed to update the weights of a neural network during training.

```
optimizer = optim.Adam(model.parameters(), lr=0.0001)
```

Additionally, for the ImprovedCNN model, dropout with a probability of 0.5 is applied before the final fully connected layer to prevent overfitting.

```
self.dropout = nn.Dropout(0.5) # Adding dropout layer
```

- **Mini-Batch Gradient Descent:**

The training process uses mini-batch gradient descent, where the model parameters are updated based on a subset of the training data.

With this method, the dataset is divided into subsets (called batches), and the gradients for each batch are computed rather than iterating over the full dataset or a single observation.

This approach enhances computational efficiency and allows iterative updates during training.

```
batch_size = 32 # Set your desired batch size
custom_training_loader =DataLoader(combined_training_data, batch_size=batch_size, shuffle=True)
custom_testing_loader=DataLoader(combined_test_data, batch_size=batch_size, shuffle=True)
custom_validation_loader=DataLoader(combined_val_data,batch_size=batch_size,shuffle=True)
```

Evaluation

1. Performance Metrics:

Model	Macro			Micro			Accuracy
	P	R	F	P	R	F	
SimpleCNN	0.49	0.47	0.415	0.468	0.483	0.475	0.48
SuperSimpleCNN	0.64	0.62	0.567	0.655	0.655	0.655	0.66
ImprovedCNN	0.62	0.61	0.55	0.638	0.638	0.638	0.62

Insights:

- The SimpleCNN model has the lowest performance in terms of macro and micro averages, as well as accuracy, among the three models.
- The SuperSimpleCNN model shows the highest precision, recall, and F1-score among all models, both in macro and micro averages. This indicates a well-balanced performance in terms of precision and recall.
- The ImprovedCNN model has a slightly lower macro average compared to SuperSimpleCNN.

Implications:

- SimpleCNN has the lowest overall performance, therefore, it may not be te best choice.
- SuperSimpleCNN demonstrates a strong balance between precision and recall, which is crucial in facial image analysis. It is the best model based on the metrics.

- ImprovedCNN offers a good balance between precision and recall, and it achieves competitive accuracy. ImprovedCNN can be considered a good choice, depending on the requirements of the application.

2. Confusion Matrix Analysis:

Here is the Confusion Matrix for each model:

SimpleCNN:

	angry	bored	focused	neutral
angry	7	2	15	13
bored	5	0	11	14
focused	0	0	33	2
neutral	0	0	12	24

SuperSimpleCNN_model:

	angry	bored	focused	neutral
angry	15	4	7	11
bored	6	6	5	13
focused	0	0	35	0
neutral	0	0	1	35

ImprovedCNN_model

	angry	bored	focused	neutral
angry	12	3	4	18
bored	6	4	5	15
focused	0	0	34	1
neutral	1	0	1	34

Frequently Confused Classes:

Model	Most confused classes	Possible Reason
SimpleCNN	<ul style="list-style-type: none">• Class 'angry' is often confused with 'focused' (13 times) and 'neutral' (10 times).• Class 'bored' is often confused with 'focused' (15 times) and 'neutral' (9 times).• Class 'focused' is most often confused with 'neutral' (14 times).	<ul style="list-style-type: none">• There might be visual similarities between facial expressions of anger and focused expressions, leading to misclassifications.• The model seems to have difficulty distinguishing between 'bored' and 'focused,' possibly due to subtle visual similarities.
SuperSimpleCNN	<ul style="list-style-type: none">• Class 'angry' is often confused with 'neutral' (13 times).• Class 'bored' is often confused with 'angry' (9 times) and 'neutral' (13 times).• Class 'focused' is often confused with 'neutral' (2 times).	<ul style="list-style-type: none">• The model shows similar challenges in distinguishing 'angry' and 'neutral' as SimpleCNN.• 'Bored' being confused with both 'angry' and 'neutral' might suggest difficulty in capturing the nuanced differences between these expressions.
ImprovedCNN	<ul style="list-style-type: none">• Class 'angry' is often confused with 'focused' (6 times) and 'neutral' (14 times).• Class 'bored' is often confused with 'neutral' (18 times).• Class 'focused' is most often confused with 'neutral' (1 time).	<ul style="list-style-type: none">• Similar to the other models, 'angry' is often confused with 'neutral,' indicating challenges in distinguishing between these two classes.• The confusion between 'bored' and 'neutral' is more pronounced in this model, suggesting a difficulty in capturing the subtle differences between these expressions.

Well Recognized Classes:

Model	Well-Recognized Class	Speculation
SimpleCNN	Class Focused is often well recognized and has a high diagonal value of 33, indicating that SimpleCNN performs well in recognizing focused expressions..	SimpleCNN might be capturing distinctive features associated with focused facial expressions. The simplicity of the model allows it to learn and generalize these features effectively.
SuperSimpleCNN	Class Angry is often well recognized and has a high diagonal value of 19, suggesting that SuperSimpleCNN is relatively successful in recognizing angry expressions.	The simplicity of SuperSimpleCNN, with fewer convolutional layers and parameters, might be beneficial for emphasizing features associated with anger, leading to better recognition.
ImprovedCNN	Class Focused is often well recognized and has a high diagonal value of 35, indicating that the ImprovedCNN model continues to perform well in recognizing focused expressions.	The improvements in the architecture of ImprovedCNN might allow it to capture more complex patterns associated with focused expressions. The added dropout layer might contribute to better generalization.

3. Impact of Architectural Variations:

The impact of depth (number of convolutional layers) on performance:

1. SimpleCNN:

- SimpleCNN has two convolutional layers.
- It may struggle to capture intricate details due to its relatively shallow architecture.
- The confusion matrix shows misclassifications, suggesting that the model might not have learned complex patterns well.

2. SuperSimpleCNN:

- SuperSimpleCNN has two convolutional layers with additional Batch Normalization and Leaky ReLU activation functions.
- The model seems to perform better than SimpleCNN, especially in correctly classifying "angry" and "bored" classes.
- Batch Normalization helps with training stability, and Leaky ReLU introduces non-linearity.

3. ImprovedCNN:

- ImprovedCNN has three convolutional layers and includes dropout for regularization.
- The model performs well, especially in distinguishing "angry" and "bored" classes.
- The deeper architecture with more convolutional layers might have helped in capturing more intricate features.

Observations:

- Deeper architectures can potentially capture more detailed features in images.
- However, increasing depth comes with a risk of overfitting, especially if the dataset is not large enough.
- Techniques like dropout and batch normalization can mitigate overfitting issues associated with deeper architectures.
- Architectural choices should be balanced based on the complexity of the task and the available dataset size.

Kernel Size Variations Impact:

The choice of kernel size in convolutional neural networks (CNNs) can significantly impact the model's ability to recognize features in images. Here's an analysis of the confusion matrices and the effects of kernel size variations on model recognition abilities:

SimpleCNN:

- Kernel Size: 5x5 in the first convolutional layer.
- Observations:
 - Struggles with recognizing certain emotions, particularly "bored" and "focused."
 - Larger kernel sizes might not capture finer details well.

SuperSimpleCNN:

- Kernel Size: 7x7 in the first convolutional layer.
- Observations:
 - Shows improvements in recognizing "angry" and "bored" emotions.
 - The larger kernel size might help capture broader facial features.

ImprovedCNN:

- Kernel Size: 3x3 in all convolutional layers.
- Observations:
 - Achieves better overall recognition performance.
 - Smaller kernel sizes might be beneficial for capturing finer details.

4. Conclusions and Forward Look:

Summary:

Based on the evaluation metrics and confusion matrices, the ImprovedCNN model performed best among the three models (SimpleCNN, SuperSimpleCNN, and ImprovedCNN). Here are the primary findings:

- Achieved the highest accuracy, precision, recall, and F1 score.
- Demonstrated better recognition across all emotion classes.
- Smaller kernel sizes (3x3) in all convolutional layers seemed effective in capturing both finer details and broader facial features.

For SimpleCNN Performance:

Struggled with recognizing certain emotions, especially "bored" and "focused."

Larger kernel size (5x5) in the first convolutional layer might not have captured finer details well.

For SuperSimpleCNN Performance:

Showed improvements in recognizing "angry" and "bored" emotions compared to SimpleCNN.

Larger kernel size (7x7) in the first convolutional layer might have helped capture broader facial features.

General Observations:

- The choice of kernel size significantly influenced recognition abilities.
- Smaller kernel sizes (ImprovedCNN) performed well in capturing finer details.
- Larger kernel sizes (SuperSimpleCNN) were effective in recognizing broader facial features.

In summary, ImprovedCNN outperformed the other models, likely due to a combination of architectural improvements, appropriate kernel sizes, and effective feature extraction. The use of smaller kernel sizes in all convolutional layers contributed to better recognition across various emotion classes.

Architecture Refinements:

By systematically exploring these refinements, it's possible to enhance the model's recognition capabilities and robustness in real-world scenarios. Here are some of the refinements:

- Experiment with Deeper Networks
- Kernel Size Optimization
- Data Augmentation
- Regularization Techniques:
- Batch Normalization and Dropout
- Learning Rate and Optimization:
- Hyperparameter Tuning

Reference Section:

<https://stackoverflow.com/questions/67595781/how-to-plot-images-in-subplots>

<https://stackoverflow.com/questions/53551410/how-to-randomly-select-images-and-put-them-to-multiple-folders-in-python>